

Solving SVP and CVP in 2^n Time Using Discrete Gaussian Sampling

Divesh Aggarwal

National University of Singapore (NUS)

Daniel Dadush

Centrum Wiskunde en Informatica (CWI)

Oded Regev


Noah Stephens-Davidowitz

New York University (NYU)

Lattices

A lattice $\mathcal{L} \subseteq \mathbb{R}^n$ is all integral combinations of some basis $B = (b_1, \dots, b_n)$.

$\mathcal{L}(B)$ denotes lattice generated by B .




Act I: The Shortest Vector Problem

Shortest Vector Problem (SVP)


Given: Lattice basis $B \in \mathbb{Q}^{n \times n}$.

Goal: Compute shortest non-zero vector in $\mathcal{L}(B)$.



Shortest Vector Problem (SVP)

$\lambda_1(\mathcal{L})$ = length of shortest non-zero vector




Algorithms for SVP

	Time	Space
[Kan86,HS07,MW15] (Enumeration)	$n^{O(n)}$	$\text{poly}(n)$
[AKS01] (Sieving)	$2^{O(n)}$	$2^{O(n)}$
[NVo8, PS09, MV10a, ...]	$2^{2.465n+o(n)}$	$2^{1.233n+o(n)}$
[MV10b] (Voronoi cell, deterministic, CVP)	$2^{2n+o(n)}$	$2^{n+o(n)}$
[ADRS15]	$2^{n+o(n)}$	$2^{n+o(n)}$

Our Algorithm

Gaussian Distribution


$$\text{Gauss}(s) := \Pr[\mathbf{x}] \propto e^{-\|\mathbf{x}\|^2/s^2}$$



$$s = 20$$

Gaussian Distribution


$$\text{Gauss}(s) := \Pr[\mathbf{x}] \propto e^{-\|\mathbf{x}\|^2/s^2}$$



$$s = 10$$

Discrete Gaussian Distribution


$$D_{\mathcal{L},s} := \Pr[\mathbf{y}] \propto e^{-\|\mathbf{y}\|^2/s^2}$$



$s = 20$

Discrete Gaussian Distribution


$$D_{\mathcal{L},s} := \Pr[\mathbf{y}] \propto e^{-\|\mathbf{y}\|^2/s^2}$$



$s = 10$

Discrete Gaussian Distribution


$$D_{\mathcal{L},s} := \Pr[\mathbf{y}] \propto e^{-\|\mathbf{y}\|^2/s^2}$$



$s = 10$

Discrete Gaussian Distribution


$$D_{\mathcal{L},s} := \Pr[\mathbf{y}] \propto e^{-\|\mathbf{y}\|^2/s^2}$$



$$s = 4$$

Discrete Gaussian Distribution


$$D_{\mathcal{L},s} := \Pr[\mathbf{y}] \propto e^{-\|\mathbf{y}\|^2/s^2}$$



$s = 4$

Discrete Gaussian Distribution


$$D_{\mathcal{L},s} := \Pr[\mathbf{y}] \propto e^{-\|\mathbf{y}\|^2/s^2}$$



$s = 2$

Discrete Gaussian Distribution


$$D_{\mathcal{L},s} := \Pr[\mathbf{y}] \propto e^{-\|\mathbf{y}\|^2/s^2}$$



$s = 2$

Discrete Gaussian Distribution


$$D_{\mathcal{L},s} := \Pr[\mathbf{y}] \propto e^{-\|\mathbf{y}\|^2/s^2}$$



$$s = 1$$

Discrete Gaussian Distribution

$$D_{\mathcal{L},s} := \Pr[\mathbf{y}] \propto e^{-\|\mathbf{y}\|^2/s^2}$$



If we can obtain “enough” samples from the discrete Gaussian with the “right” (small) parameter, then we can solve SVP.

$$s = 2$$

Discrete Gaussian Distribution

We need at most 1.38^n vectors with $s \approx \lambda_1(\mathcal{L})/\sqrt{n}$ [KL78].
(uses bounds on the kissing number)

$D_{\mathcal{L},s}$ is very well-studied for very high parameters, $s \gtrsim \lambda_n(\mathcal{L})$, i.e. above the “smoothing parameter” of the lattice.


[Kle00, GPVo8] show how to sample in this regime in polynomial time.

(Previously could not do much better, even in exponential time.)

Discrete Gaussian Distribution

Easy

[Kle00, GPVo8]




$$s \gg \lambda_1(\mathcal{L})$$

Our goal



Hard




$$s \approx \lambda_1(\mathcal{L})/\sqrt{n}$$


Can we use samples from the LHS to get samples from the RHS?

Discrete Gaussian Distribution

$$\mathbf{x} \sim \text{Gauss}(s)$$



$$\frac{\mathbf{x}}{2} \sim \text{Gauss}(s/2)$$




=


2

Discrete Gaussian Distribution


$$\mathbf{y} \sim D_{\mathcal{L},s}$$



||



Discrete Gaussian Distribution



Converting Gaussian Vectors

What if we *condition* on the result being in the lattice?

$$\Pr_{\mathbf{y} \sim D_{\mathcal{L},s}} \left[\frac{\mathbf{y}}{2} = \mathbf{x} \mid \frac{\mathbf{y}}{2} \in \mathcal{L} \right] \propto e^{-4\|\mathbf{x}\|^2/s^2}$$


↑
Progress!

Unfortunately, this requires us to throw out a lot of vectors.

We only keep one from every $\approx 2^n$ vectors each time we do this, leading to a very slow algorithm!


Converting Gaussian Vectors

$$\mathbf{x}_1 \sim \text{Gauss}(s)$$




+

$$\mathbf{x}_2 \sim \text{Gauss}(s)$$



=


$$\frac{\mathbf{x}_1 + \mathbf{x}_2}{2} \sim \text{Gauss}(s/\sqrt{2})$$



2


Converting Gaussian Vectors

$$\mathbf{y}_1 \sim D_{\mathcal{L},s}$$




+

$$\mathbf{y}_2 \sim D_{\mathcal{L},s}$$




$\stackrel{?}{=}$

$$\frac{\mathbf{y}_1 + \mathbf{y}_2}{2} \sim D_{\mathcal{L},s/\sqrt{2}}$$



2

Converting Gaussian Vectors



Converting Gaussian Vectors

What about the average of two discrete Gaussian vectors *conditioned on* the result being in the lattice?

Converting Gaussian Vectors

When do we have $\frac{\mathbf{y}_1 + \mathbf{y}_2}{2} \in \mathcal{L}$?

$$\mathbf{y}_1 = a_1 \mathbf{b}_1 + \dots + a_n \mathbf{b}_n \quad \mathbf{y}_2 = a_2 \mathbf{b}_1 + \dots + a_{2,n} \mathbf{b}_n$$

We have $(\mathbf{y}_1 + \mathbf{y}_2)/2 \in \mathcal{L}$ if and only if
 $\mathbf{y}_1, \mathbf{y}_2$ are in the same **coset** of $2\mathcal{L}$.
 (Note that there are 2^n cosets)

$$\frac{\mathbf{y}_1 + \mathbf{y}_2}{2} \in \mathcal{L} \iff a_{1,i} \equiv a_{2,i} \pmod{2}$$


$$\iff \mathbf{y}_1 \equiv \mathbf{y}_2 \pmod{2\mathcal{L}}$$

Converting Gaussian Vectors

What about the average of two discrete Gaussian vectors *conditioned on* the result being in the lattice?

Converting Gaussian Vectors


$$\mathcal{L}^\dagger = \{(y_1, y_2) \in \mathcal{L} : y_1 \equiv y_2 \pmod{2\mathcal{L}}\}$$



What about the average of two discrete Gaussian vectors *conditioned* on the result being in the lattice?

Converting Gaussian Vectors


$$\text{avg}(y_1, y_2) = \left(\frac{y_1 + y_2}{2}, \frac{y_1 - y_2}{2} \right)$$



What about the average of two discrete Gaussian vectors *conditioned* on the result being in the lattice?

Converting Gaussian Vectors

$$\text{avg}(\mathcal{L}^\dagger) = \mathcal{L} \times \mathcal{L}$$



What about the average of two discrete Gaussian vectors *conditioned* on the result being in the lattice?

Converting Gaussian Vectors

$$\text{avg}(\mathcal{L}^\dagger) = \mathcal{L} \times \mathcal{L}$$

$\| \cdot \|_2$

If we sample $\mathbf{y}_1, \mathbf{y}_2 \sim D_{\mathcal{L},s}$,
then their average will be distributed as $D_{\mathcal{L},s/\sqrt{2}}$,
if we condition on the result being in the lattice.

$\sqrt{2}$

$$(\mathbf{y}_1, \mathbf{y}_2) \sim D_{\mathcal{L}^\dagger,s} \Rightarrow \text{avg}(\mathbf{y}_1, \mathbf{y}_2) \sim D_{\mathcal{L} \times \mathcal{L},s/\sqrt{2}}$$



$$D_{\mathcal{L},s} := \Pr[\mathbf{y}] \propto e^{-\|\mathbf{y}\|^2/s^2}$$

Progress!

Stitching a Discrete Gaussian Together

$$\Pr_{y_1, y_2 \sim D_{\mathcal{L}, s}} \left[\frac{y_1 + y_2}{2} = y \mid \frac{y_1 + y_2}{2} \in \mathcal{L} \right] \\ \propto \sum_{c \in \mathcal{L} \pmod{2\mathcal{L}}} \Pr[D_{\mathcal{L}, s} \in c]^2 \Pr_{y_1, y_2 \sim D_{2\mathcal{L}+c, s}} \left[\frac{y_1 + y_2}{2} = y \right]$$

Generating a single $D_{\mathcal{L}, s/\sqrt{2}}$ sample:

1. Sample $c \in \mathcal{L} \pmod{2\mathcal{L}}$ with probability $\propto \Pr[D_{\mathcal{L}, s} \in c]^2$.
2. Output $(Y_1 + Y_2)/2$ where $Y_1, Y_2 \sim D_{2\mathcal{L}+c, s}$.

Discrete Gaussian Combiner

Input: Y_1, \dots, Y_M iid $D_{\mathcal{L},s}$ samples ($M \approx 2^n$)

1. “Bucket” samples according to their coset (mod $2\mathcal{L}$).

2. Repeat many times:

1. Sample coset \mathbf{c} with probability $\propto \Pr[D_{\mathcal{L},s} \in \mathbf{c}]^2$.

2. Output $(Y_i + Y_j)/2$, for $Y_i, Y_j \in \mathbf{c}$.

3. Remove Y_i, Y_j from list.



Don't have access to
this distribution!

Rejection Sampling

Achieving $\propto \Pr[D_{\mathcal{L},s} \in \mathbf{c}]^2$:

Same as trivial strategy!

First Pass:

Sample $\mathbf{c} \sim D_{\mathcal{L},s} \pmod{2\mathcal{L}}$.

Accept \mathbf{c} with probability $\Pr[D_{\mathcal{L},s} \in \mathbf{c}]$ o/w reject.



Implementation:

Sample $Y_1 \sim D_{\mathcal{L},s}$ and let \mathbf{c} be $Y_1 \pmod{2\mathcal{L}}$.

Sample $Y_2 \sim D_{\mathcal{L},s}$.

Output \mathbf{c} if $Y_1 \equiv Y_2 \pmod{2\mathcal{L}}$.

Rejection Sampling

Achieving $\propto \Pr[D_{\mathcal{L},s} \in c]^2$:

Second Try:

Sample $c \sim D_{\mathcal{L},s} \pmod{2\mathcal{L}}$.

Accept c with probability $\frac{\Pr[D_{\mathcal{L},s} \in c]}{p_{\max}}$ o/w reject,

where

$$p_{\max} = \max_{b \in \mathcal{L} \pmod{2\mathcal{L}}} \Pr[D_{\mathcal{L},s} \in b]$$

Implementation: ???

Discrete Gaussian Combiner

Input: Y_1, \dots, Y_M iid $D_{\mathcal{L},s}$ samples ($M \approx 2^n$)

Use first $M/6$ samples to estimate p_{\max} .

	1	...	$Mp_{\max}/3$
$\mathcal{L}(\text{mod } 2\mathcal{L})$ 2^n buckets			
	# samples in		
	each bucket		
	First $1/p_{\max}$ samples	...	Last $1/p_{\max}$ samples

Discrete Gaussian Combiner

Input: Y_1, \dots, Y_M iid $D_{\mathcal{L},s}$ samples ($M \approx 2^n$)

1. Compute p_{\max} and bucket counts (previous slide).
2. For i ranging over last $M/6$ samples:
 1. Let $\mathbf{c} = Y_i \pmod{2\mathcal{L}}$.
 2. Find first unused bucket count $k_{\mathbf{c}}$ for coset \mathbf{c} .
 3. With probability $\min \{1, k_{\mathbf{c}}/n^{O(1)}\}$,
output $(Y_i + Y_j)/2$
where Y_j is any sample contributing to $k_{\mathbf{c}}$.

How Many Vectors Do We Get?

$M := \# \text{ input vectors}$

$$\# \text{ output vectors} \approx M \cdot \frac{\sum_c \Pr[D_{\mathcal{L},s} \in c]^2}{\max_b \Pr[D_{\mathcal{L},s} \in b]}$$

Worst case bound: probability is at least $1/\sqrt{|\text{support}|}$.

Max pop to $1/2^n$ after n steps!

How Many Vectors Do We Get?

$$\# \text{ of output vectors} \approx M \cdot \frac{\sum_{\mathbf{c}} \Pr[D_{\mathcal{L},s} \in \mathbf{c}]^2}{\max_{\mathbf{c}} \Pr[D_{\mathcal{L},s} \in \mathbf{c}]}$$

$$\begin{aligned} \sum_{\mathbf{c}} \Pr[D_{\mathcal{L},s} \in \mathbf{c}]^2 &= \frac{\rho_s(\mathcal{L}^\dagger)}{\rho_s(\mathcal{L})^2} \\ &= \frac{\rho_s(\sqrt{2}\mathcal{L} \times \sqrt{2}\mathcal{L})}{\rho_s(\mathcal{L})^2} \\ &= \frac{\rho_{s/\sqrt{2}}(\mathcal{L})^2}{\rho_s(\mathcal{L})^2} \end{aligned}$$

$$\rho_s(\mathcal{L}) := \sum_{\mathbf{y} \in \mathcal{L}} e^{-\|\mathbf{y}/s\|^2}$$

How Many Vectors Do We Get?

$$\begin{aligned}
 \# \text{ of output vectors} &\approx M \cdot \frac{\rho_{s/\sqrt{2}}(\mathcal{L})^2}{\rho_s(\mathcal{L})^2 \max_{\mathbf{c}} \Pr[D_{\mathcal{L},s} \in \mathbf{c}]} \\
 \max_{\mathbf{c}} \rho_s(2\mathcal{L} + \mathbf{c}) &\xrightarrow{\quad} = M \cdot \frac{\rho_{s/\sqrt{2}}(\mathcal{L})^2}{\rho_s(\mathcal{L}) \rho_s(2\mathcal{L})} \\
 = \rho_s(2\mathcal{L}) & \\
 &= M \cdot \frac{\rho_{s/\sqrt{2}}(\mathcal{L})^2}{\rho_s(\mathcal{L}) \rho_{s/2}(\mathcal{L})}
 \end{aligned}$$

$$\rho_s(\mathcal{L}) := \sum_{\mathbf{y} \in \mathcal{L}} e^{-\|\mathbf{y}/s\|^2}$$

How Many Vectors Do We Get?

$$\# \text{ of output vectors after } \ell \text{ steps} \approx M \cdot \prod_{i=0}^{\ell} \frac{\rho_{2^{-\frac{i+1}{2}}s}(\mathcal{L})^2}{\rho_{2^{-\frac{i}{2}}s}(\mathcal{L})\rho_{2^{-\frac{i+2}{2}}s}(\mathcal{L})}$$

Recall that we only need 1.38^n samples to solve SVP!

$$\rho_s(\mathcal{L}) \leq 2^{n/2} \rho_{s/\sqrt{2}}(\mathcal{L}) \longrightarrow \geq M \cdot 2^{-n/2}$$

Setting $M \approx 2^n$ gives $\# \text{ output vectors} \approx 2^{n/2}$

Key Estimates

Poisson summation formula: “nice” function f

$$\sum_{\mathbf{y} \in \mathcal{L}} f(\mathbf{y} + \mathbf{t}) = \frac{1}{\det(\mathcal{L})} \sum_{\mathbf{x} \in \mathcal{L}^*} \hat{f}(\mathbf{x}) e^{2\pi i \langle \mathbf{x}, \mathbf{t} \rangle}$$

Plug in $e^{-\pi \|x/s\|^2}$:

$$\rho_s(\mathcal{L} + \mathbf{t}) = \frac{s^n}{\det(\mathcal{L})} \sum_{\mathbf{x} \in \mathcal{L}^*} e^{-\pi \|s\mathbf{x}\|^2} e^{2\pi i \langle \mathbf{x}, \mathbf{t} \rangle}$$

$$\rho_s(\mathcal{L}) = \frac{s^n}{\det(\mathcal{L})} \rho_{1/s}(\mathcal{L}^*)$$

Key Estimates

$$\rho_s(\mathcal{L} + \mathbf{t}) = \frac{s^n}{\det(\mathcal{L})} \sum_{\mathbf{x} \in \mathcal{L}^*} e^{-\pi \|s\mathbf{x}\|^2} e^{2\pi i \langle \mathbf{x}, \mathbf{t} \rangle}$$

$$\rho_s(\mathcal{L}) = \frac{s^n}{\det(\mathcal{L})} \rho_{1/s}(\mathcal{L}^*)$$

Corollary 1: $\max_{\mathbf{t}} \rho_s(\mathcal{L} + \mathbf{t}) = \rho_s(\mathcal{L})$

Corollary 2: $\rho_{\alpha s}(\mathcal{L}) \leq \alpha^n \rho_s(\mathcal{L})$ for $\alpha \geq 1$.

Final Algorithm

SVPSolver(\mathcal{L})


1. Use GPV to get $\approx 2^n$ samples from $D_{\mathcal{L},s}$ with $s \gg \lambda_1(\mathcal{L})$.
2. Run the (“squaring”) discrete Gaussian combiner on the result repeatedly.
3. Output $\approx 2^{n/2}$ samples from $D_{\mathcal{L},s}$ with $s \approx \lambda_1(\mathcal{L})/\sqrt{n}$.
4. We can then simply output a shortest non-zero vector from our samples.

Act II: The Closest Vector Problem

Closest Vector Problem (CVP)

Given: Lattice basis $B \in \mathbb{Q}^{n \times n}$, target $\mathbf{t} \in \mathbb{Q}^n$.

Goal: Compute $\mathbf{y} \in \mathcal{L}(B)$ minimizing $\|\mathbf{t} - \mathbf{y}\|$.



Closest Vector Problem (CVP)

CVP seems to be the harder problem:
there is a dimension preserving reduction
from **SVP** to **CVP** [GMSS99].

Algorithms for CVP


	Time	CVP?	Deterministic?
[Kan86, HSo7, MW15] (Enumeration)	$n^{O(n)}$	Yes	Yes
[AKSo2, BNo9, HPS11, ...] (Sieving)	$2^{O(n)}$	Approximate	No
[MV10b] (Voronoi cell)	$2^{2n+o(n)}$	Yes	Yes
[ADRS15] (Discrete Gaussian)	$2^{n+o(n)}$	Approximate	No
[ADS15]	$2^{n+o(n)}$	Yes	No

Disclaimer

The algorithm is quite complicated, so the following is a over-simplified high level sketch.

The Discrete Gaussian Distribution

$$D_{\mathcal{L}, \mathbf{t}, s} := \Pr[\mathbf{y}] \propto e^{-\|\mathbf{y} - \mathbf{t}\|^2 / s^2}$$




$$s = 20$$

$$\mathcal{L} := \mathbb{Z}^2$$

$$\mathbf{t} := (0, 5/2)$$

The Discrete Gaussian Distribution

$$D_{\mathcal{L}, \mathbf{t}, s} := \Pr[\mathbf{y}] \propto e^{-\|\mathbf{y} - \mathbf{t}\|^2 / s^2}$$




$$s = 10$$

$$\mathcal{L} := \mathbb{Z}^2$$

$$\mathbf{t} := (0, 5/2)$$

The Discrete Gaussian Distribution

$$D_{\mathcal{L}, \mathbf{t}, s} := \Pr[\mathbf{y}] \propto e^{-\|\mathbf{y} - \mathbf{t}\|^2 / s^2}$$




$$s = 10$$

$$\mathcal{L} := \mathbb{Z}^2$$

$$\mathbf{t} := (0, 5/2)$$

The Discrete Gaussian Distribution

$$D_{\mathcal{L}, \mathbf{t}, s} := \Pr[\mathbf{y}] \propto e^{-\|\mathbf{y} - \mathbf{t}\|^2 / s^2}$$




$$s = 5$$

$$\mathcal{L} := \mathbb{Z}^2$$

$$\mathbf{t} := (0, 5/2)$$

The Discrete Gaussian Distribution

$$D_{\mathcal{L}, \mathbf{t}, s} := \Pr[\mathbf{y}] \propto e^{-\|\mathbf{y} - \mathbf{t}\|^2 / s^2}$$




$$s = 5$$

$$\mathcal{L} := \mathbb{Z}^2$$

$$\mathbf{t} := (0, 5/2)$$

The Discrete Gaussian Distribution

$$D_{\mathcal{L}, \mathbf{t}, s} := \Pr[\mathbf{y}] \propto e^{-\|\mathbf{y} - \mathbf{t}\|^2 / s^2}$$




$$s = 2$$

$$\mathcal{L} := \mathbb{Z}^2$$

$$\mathbf{t} := (0, 5/2)$$

The Discrete Gaussian Distribution

$$D_{\mathcal{L}, \mathbf{t}, s} := \Pr[\mathbf{y}] \propto e^{-\|\mathbf{y} - \mathbf{t}\|^2 / s^2}$$




$$s = 2$$

$$\mathcal{L} := \mathbb{Z}^2$$

$$\mathbf{t} := (0, 5/2)$$

The Discrete Gaussian Distribution

$$D_{\mathcal{L}, \mathbf{t}, s} := \Pr[\mathbf{y}] \propto e^{-\|\mathbf{y} - \mathbf{t}\|^2 / s^2}$$



$$s = 1$$


$$\mathcal{L} := \mathbb{Z}^2$$

$$\mathbf{t} := (0, 5/2)$$

The Discrete Gaussian Distribution

$$D_{\mathcal{L}, \mathbf{t}, s} := \Pr[\mathbf{y}] \propto e^{-\|\mathbf{y} - \mathbf{t}\|^2 / s^2}$$

CVP trivially reduces to sampling from the discrete Gaussian distribution $D_{\mathcal{L}, \mathbf{t}, s}$ for a small enough parameter s .



$$s = 1$$

$$\mathcal{L} := \mathbb{Z}^2$$

$$\mathbf{t} := (0, 5/2)$$

“Rotation” Identity Generalizes

$$\Pr_{\mathbf{y}_1, \mathbf{y}_2 \sim D_{\mathcal{L}, s}} \left[\frac{\mathbf{y}_1 + \mathbf{y}_2}{2} = \mathbf{y} \mid \frac{\mathbf{y}_1 + \mathbf{y}_2}{2} \in \mathcal{L} \right] = \Pr_{\mathbf{x} \sim D_{\mathcal{L}, s/\sqrt{2}}} [\mathbf{x} = \mathbf{y}]$$


Great! So, we just need to run the squaring combiner and we’re done! Right!?

$$\Pr_{\mathbf{y}_1, \mathbf{y}_2 \sim D_{\mathcal{L}, t, s}} \left[\frac{\mathbf{y}_1 + \mathbf{y}_2}{2} = \mathbf{y} \mid \frac{\mathbf{y}_1 + \mathbf{y}_2}{2} \in \mathcal{L} \right] = \Pr_{\mathbf{x} \sim D_{\mathcal{L}, t, s/\sqrt{2}}} [\mathbf{x} = \mathbf{y}]$$

Initialization Issues

- The [GPVo8] sampler does work for sampling shifted $D_{\mathcal{L}, \mathbf{t}, s}$, but gives
- When s is too small, the sampler is not effective.
- When $\mathbf{t} \neq 0$, we may not be able to do this.
- So, we must initialize with $s \gtrsim \text{dist}(\mathbf{t}, \mathcal{L})$.

Even if apply the combiner n times, we can only sample at $s \approx 2^{-n} \text{dist}(\mathbf{t}, \mathcal{L})$.



Combiner Loss Factor

Going from $s \rightarrow s/\sqrt{2}$:

Center

No obvious “magical cancelation”.

General \mathbf{t} :

$$\frac{\rho_{s/\sqrt{2}}(\mathcal{L})\rho_{s/\sqrt{2}}(\mathcal{L}-\mathbf{t})}{\rho_s(\mathcal{L}) \max_{\mathbf{c} \in \mathcal{L}/2\mathcal{L}} \rho_s(\mathbf{c}-\mathbf{t})}$$

Combiner Loss Factor

Theorem: Combiner loss going from


$$s \rightarrow s_k := s2^{-k/2}$$

is no worse than

$$\frac{2^{-n}}{\max_{c \in \mathcal{L}/2\mathcal{L}} \Pr[D_{\mathcal{L}, t, s_k} \in c]}.$$

If we start with $2^{n+o(n)}$ samples, we always “see” the heaviest coset at each stage.


Exact vs Approximate CVP



We Need Small Parameters

The reduction from CVP to DGS needs $s \ll \lambda_1(\mathcal{L})$, but we can only handle $s \approx 2^{-n} \cdot \text{dist}(\mathbf{t}, \mathcal{L})$.


For such parameters, we obtain approximate solutions with unreasonably good approximation factor $\gamma \approx 1 + 2^{-n}$, but not **exact** solutions.



Hope for exact CVP

To apply recursion, need to identify them and show that there are not too many.


$2^{n+o(n)}$ time \coloneqq at most 2 sub-problems per dimension!



Clusters

Claim: There are at most 2^n exact closest vectors.


Must lie in different cosets of $\mathcal{L}/2\mathcal{L}$.



Sphere containing the
closest vectors


Clusters

Claim: The approximate closest vectors are contained in 2^n “clusters” of small diameter.



Clusters

Claim: $\sqrt{1 + \epsilon^2}$ approx. CVP sols u and v .
 $v - u \in 2\mathcal{L}$ implies $\|v - u\| \leq 2\epsilon \cdot \text{dist}(\mathbf{t}, \mathcal{L})$.



Clusters

Claim: $\sqrt{1 + \epsilon^2}$ approx. CVP sols u and v .
 $v - u \in 2\mathcal{L}$ implies $\|v - u\| \leq 2\epsilon \cdot \text{dist}(\mathbf{t}, \mathcal{L})$.

$$\begin{aligned}\|v - u\|^2 &= 2\|v - \mathbf{t}\|^2 + 2\|u - \mathbf{t}\|^2 \\ &\quad - 4\|(v + u)/2 - \mathbf{t}\|^2 \\ &\leq 4(1 + \epsilon^2) \cdot \text{dist}(\mathbf{t}, \mathcal{L})^2 \\ &\quad - 4 \cdot \text{dist}(\mathbf{t}, \mathcal{L})^2 \\ &= 4\epsilon^2 \cdot \text{dist}(\mathbf{t}, \mathcal{L})^2\end{aligned}$$

Taking advantage of clusters

“nearly orthogonal” basis b_1, \dots, b_n of \mathcal{L}
(lengths in approx. non-decreasing order)

$1 + 2^{-n}$ approx CVP sols y_1, \dots, y_N for \mathbf{t} .

$$y_j = \sum_i a_{i,j} b_i \quad \forall j$$

Theorem: $\exists k$ such that last k coefficients

$$\{(a_{n-k+1,j}, \dots, a_{n,j}) : j \in [N]\}$$

come from set of size $\approx 2^k$.

Recurse on these!

Taking advantage of clusters

Assume: **orthogonal** lattice \mathcal{L}

$$\mathcal{L} = \{(x_1 b_1, \dots, x_n b_n) : x \in \mathbb{Z}^n\} \\ (0 \leq b_1 \leq \dots \leq b_n)$$

For $\epsilon = 2^{-n}$, **all** coordinates are fixed by parity unless there are **exponential** gaps in basis vector lengths.
But such gaps can exist....

Claim: If $y_r - y_s \in 2\mathcal{L}$ and $b_{n-k+1} > \sqrt{n} \epsilon b_n$

then $(a_{n-k+1,r}, \dots, a_{n,r}) = (a_{n-k+1,s}, \dots, a_{n,s})$

Taking advantage of clusters

Claim: If $y_r - y_s \in 2\mathcal{L}$ and $b_{n-k+1} \geq \sqrt{n} \epsilon b_n$

$$1. \text{dist}(\mathbf{t}, \mathcal{L}) \leq \frac{1}{2} \sqrt{\sum_i b_i^2} \leq \frac{\sqrt{n}}{2} b_n$$

This shows we have at most 2^n clusters each of which is $n - k$ dimensional, but we need 2^k clusters!!!

$-k+1$


If y_r, y_s differ on any coordinate

$$i \in \{n - k + 1, \dots, n\}$$

their difference would have norm at least b_{n-k+1} .


Exploiting gaps in basis lengths

Idea: Only match parity on “high order bits”.



Exploiting gaps in basis lengths


Idea: Only match parity on “high order bits”.



Parity of last k coefficients determines these coefficients exactly.


Exploiting gaps in basis lengths

Idea: Only match parity on “high order bits”.



Idea: Can round first $n - k$ coefficients to desired parity without increasing distance to \mathbf{t} by much.

Exploiting gaps in basis lengths




What if there are no large gaps?

Exploiting gaps in basis lengths

Idea: Again only match parity on last k bits.

Can guarantee k is large in this case.



High Level Algorithm

Input: n -dimensional lattice \mathcal{L} and target \mathbf{t} .

Output: Closest lattice vectors in \mathcal{L} to \mathbf{t} .

1. Compute **short** basis B of \mathcal{L} , and number k of “**high order coordinates**”.
2. Get many $1 + 2^{-n}$ approx. closest vectors via **DGS**.
3. Group them according to last k coordinates with respect to B and recurse on each group .

Complexity Sketch

Initialization: (one shot $2^{n+o(n)}$ time)

Compute **short** basis B of \mathcal{L} , and number k of “**high order coordinates**” (can compute for each rec. level).

Per level work: ($2^{n+o(n)}$ time)

Sample many approx. closest vectors via **DGS**.

Recursion: ($\approx 2^k$ subproblems of dim. $n - k$)

Group them according to last k coordinates with respect to B and recurse.

Total runtime: $2^{n+o(n)}$

Key Challenges

Runtime:

1. Getting many **DGS** samples at low parameters.
2. Show last **k coeffs** \approx determined by their parity.
3. Deal with $\approx 2^k$ subproblems in recursion analysis.

Correctness:

Show that we **hit last k coeffs** of an exact closest vector with high probability.



Summary of Results

Discussed in this talk

- $2^{n+o(n)}$ algorithm for SVP and CVP.
- How to sample $2^{n/2}$ vectors from $D_{\mathcal{L},s}$ for any s in time $2^{n+o(n)}$

Additional results from this work


- $2^{n/2+o(n)}$ -time algorithm for sampling $2^{n/2}$ vectors above smoothing.
- 1.93-GapSVP.
- .422-BDD.

Recent work

- Sampling from DGS reduces to SVP. [Ste16]
(not equivalence because the reduction in the other direction requires $1.38^n D_{\mathcal{L},s}$ samples.)

Open Questions/Future Work

- Other uses for discrete Gaussian sampling at arbitrary parameters?
- Faster discrete centered Gaussian sampling?
- Strong lower bounds for CVP/SVP assuming SETH (or something similar)?
- Deterministic / Las Vegas algorithms with same complexity?



Thanks!

